

# CS294: Deep Reinforcement Learning

## Homework 2 - Policy Gradients

Siyao Jia, SID: 3032531126

### Problem 1. State-dependent baseline

(a)

$$\begin{aligned} & \sum_{t=1}^T E_{\tau \sim p_{\theta}(\tau)} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) b(s_t)] \\ &= \sum_{t=1}^T E_{s_t, a_t \sim p_{\theta}(s_t, a_t)} [E_{p_{\theta}(\tau | s_t, a_t | s_t, a_t)} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) b(s_t)]] \\ &= \sum_{t=1}^T E_{s_t, a_t \sim p_{\theta}(s_t, a_t)} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) b(s_t)] \\ &= \sum_{t=1}^T b(s_t) E_{s_t, a_t \sim p_{\theta}(s_t, a_t)} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)] \\ &= \sum_{t=1}^T b(s_t) E_{s_t} [E_{a_t \sim \pi_{\theta}(a_t | s_t)} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]] \\ &= \sum_{t=1}^T b(s_t) E_{s_t} \left[ \int \pi_{\theta}(a_t | s_t) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) da_t \right] \\ &= \sum_{t=1}^T b(s_t) E_{s_t} \left[ \int \pi_{\theta}(a_t | s_t) \frac{\nabla_{\theta} \pi_{\theta}(a_t | s_t)}{\pi_{\theta}(a_t | s_t)} da_t \right] \\ &= \sum_{t=1}^T b(s_t) E_{s_t} [\nabla_{\theta} \int \pi_{\theta}(a_t | s_t) da_t] \\ &= 0 \end{aligned} \tag{1}$$

(b)

(a) Because in Markov process, the current state  $s_{t+1}$  only depends on previous  $s_t$  and  $a_t$ , and is independent of  $s_{t-1}$  or  $a_{t-1}$ . Therefore conditioning on  $(s_1, a_1, \dots, a_{t^*-1}, s_{t^*})$  is equivalent to conditioning only on  $s_{t^*}$ .

(b)

$$\begin{aligned} & \sum_{t=1}^T E_{\tau \sim p_{\theta}(\tau)} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) b(s_t)] \\ &= \sum_{t=1}^T E_{p_{\theta}(s_{1:t}, a_{1:t-1})} [E_{p_{\theta}(s_{t+1:T}, a_{t:T} | s_{1:t}, a_{1:t-1})} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) b(s_t)]] \\ &= \sum_{t=1}^T E_{p_{\theta}(s_{1:t}, a_{1:t-1})} [E_{p_{\theta}(s_{t+1:T}, a_{t:T} | s_t)} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) b(s_t)]] \\ &= \sum_{t=1}^T E_{p_{\theta}(s_{1:t}, a_{1:t-1})} [b(s_t) E_{p_{\theta}(s_{t+1:T}, a_{t:T} | s_t)} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]] \\ &= \sum_{t=1}^T E_{p_{\theta}(s_{1:t}, a_{1:t-1})} [b(s_t) E_{\pi_{\theta}(a_t | s_t)} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]] \\ &= 0 \end{aligned} \tag{2}$$

### Problem 2. Neural networks

See details in code.

### Problem 3. Policy Gradient

See details in code.

### Problem 4. CartPole

- Figure 1 compares the learning curves for the experiments prefixed with sb\_ (The small batch experiments). Figure 2 compare the learning curves for the experiments prefixed with lb\_. (The large batch experiments.)
- The gradient estimator using reward-to-go has better performance without advantage-centering.
  - Yes, advantage centering helps slightly.
  - Yes, batch size makes a huge impact by stabilizing the return value.
- The command line configurations I used:
  - "python train\_pg\_f18.py CartPole-v0 -n 100 -b 1000 -e 3 -dna --exp\_name sb\_no\_rtg\_dna"
  - "python train\_pg\_f18.py CartPole-v0 -n 100 -b 1000 -e 3 -rtg -dna --exp\_name sb\_rtg\_dna"
  - "python train\_pg\_f18.py CartPole-v0 -n 100 -b 1000 -e 3 -rtg --exp\_name sb\_no\_rtg\_na"
  - "python train\_pg\_f18.py CartPole-v0 -n 100 -b 5000 -e 3 -dna --exp\_name lb\_no\_rtg\_dna"
  - "python train\_pg\_f18.py CartPole-v0 -n 100 -b 5000 -e 3 -rtg -dna --exp\_name sb\_rtg\_dna"
  - "python train\_pg\_f18.py CartPole-v0 -n 100 -b 5000 -e 3 -rtg --exp\_name sb\_no\_rtg\_na"




Figure 1: Learning Curves for Small Batch Experiments




Figure 2: Learning Curves for Large Batch Experiments




Figure 3: Learning Curve for InvertedPendulum

### Problem 5. IvertedPendulum

1. I choose batch size  $b$  of 5000 and learning rate  $lr$  of 0.005. The learning curve is showed in Figure 3.
2. The command line configurations I used:  
`"python train_pg_f18.py InvertedPendulum-v2 -ep 1000 -discount 0.9 -n 100 -e 3 -l 2 -s 64 -b 5000 -lr 0.005 -rtg -exp_name hc_b5000_r0.005"`

### Problem 6. Neural network baseline

See details in code.

### Problem 7. LunarLander

The learning curve from the following command is plotted in Figure 4

`"python train_pg_f18.py LunarLanderContinuous-v2 -ep 1000 -discount 0.99 -n 100 -e 3 -l 2 -s 64 -b 40000 -lr 0.005 -rtg -nn_baseline -exp_name ll_b40000_r0.005"`

### Problem 8. HalfCheetah

1. Larger batch size reduces variance and stabilizes the performance. Lower learning rate takes longer time, but are more likely to converge. Vice versa for high learning rate.
2. I choose batch size  $b$  of 50000 and learning rate  $lr$  of 0.02. The learning curves of the




Figure 4: Learning Curve for LunarLander




Figure 5: Learning Curves for HalfCheetah

following commands is plotted in Figure 5

```
"python train_pg_f18.py HalfCheetah-v2 -ep 150 -discount 0.95 -n 100 -e 3 -l 2 -s 32 -b 50000 -lr 0.02 -exp_name hc_b50000_r0.02"
```

```
"python train_pg_f18.py HalfCheetah-v2 -ep 150 -discount 0.95 -n 100 -e 3 -l 2 -s 32 -b 50000 -lr 0.02 -rtg -exp_name hc_rtg_b50000_r0.02"
```

```
"python train_pg_f18.py HalfCheetah-v2 -ep 150 -discount 0.95 -n 100 -e 3 -l 2 -s 32 -b 50000 -lr 0.02 -nn_baseline -exp_name hc_baseline_b50000_r0.02"
```

```
"python train_pg_f18.py HalfCheetah-v2 -ep 150 -discount 0.95 -n 100 -e 3 -l 2 -s 32 -b 50000 -lr 0.02 -rtg -nn_baseline -exp_rtg_baseline_name hc_b50000_r0.02"
```